

AI & Applications

Lecturer: Hoang Duc Quy

Email: quyhd@ut.edu.vn

Department: Institute of Information Technology and
Electrical- Electronic Engineering (IT3E)

MACHINE LEARNING ALGORITHMs

Linear Regression

What you'll learn?

- Linear regression model
- Cost function
 - Cost function formula
 - Cost function intuition
 - Visualize the cost function
 - Example
- Gradient descent algorithm*
 - Theory and implementation
 - Learning rate
 - GD for linear regression

Terminology

Training set: **data used to train the model**

x = “input” variable

Size in feet²

Price in \$1000

or **feature**

600

150

y = “output” variable

615

210

or **target**

800

250

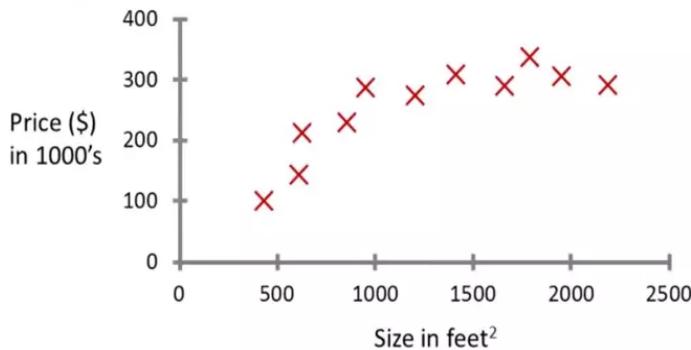
m = number of training examples

...

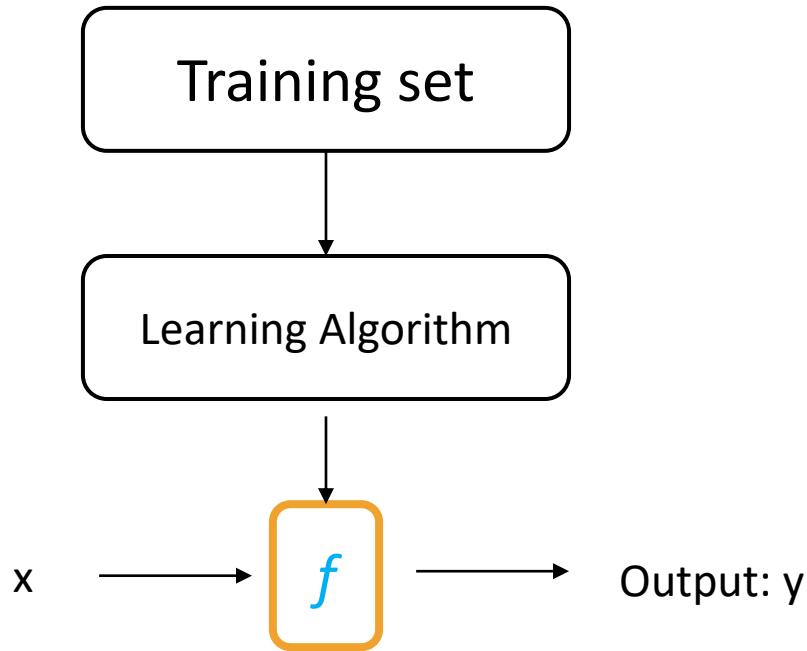
...

Housing price prediction.

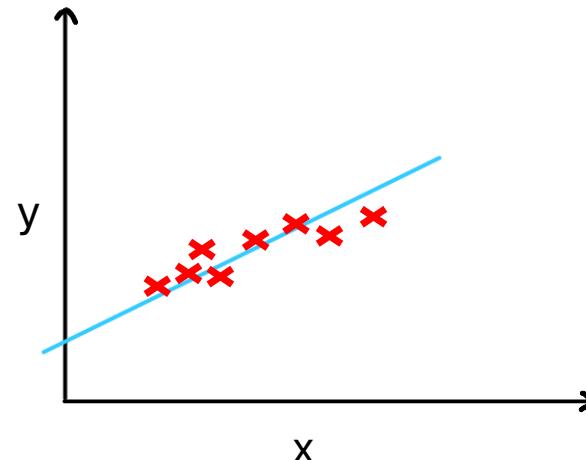
$(x^{(i)}, y^{(i)})$ = i^{th} training example



Linear Regression model



How to represent f ?



$$\text{Model: } f_{w,b}(x) = wx + b$$

- w is the weight
- b is the bias
- $f_{w,b}(x)$ is the prediction
- Together w and b are called parameters

Training set: **data used to train** the model

Size in feet²

600

615

800

...

Price in \$1000

150

210

250

...

x = “input” variable
or **feature**

y = “output” variable

or **target**

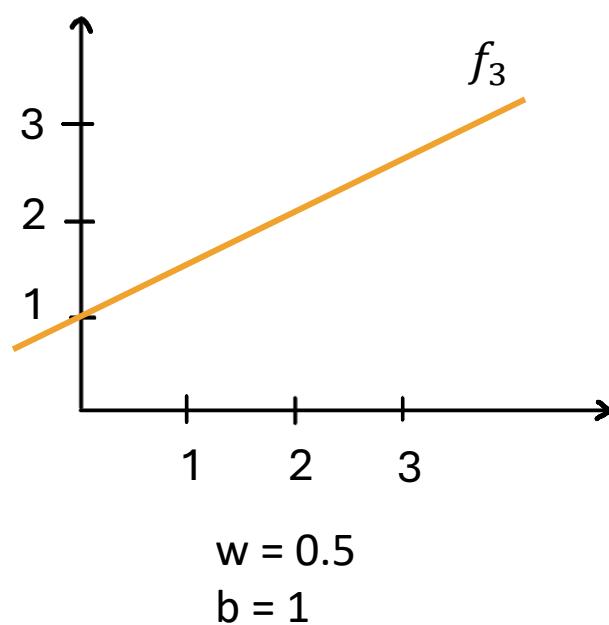
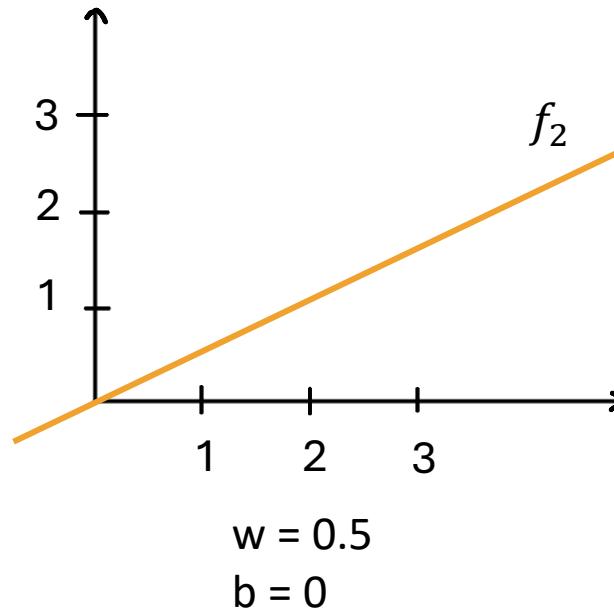
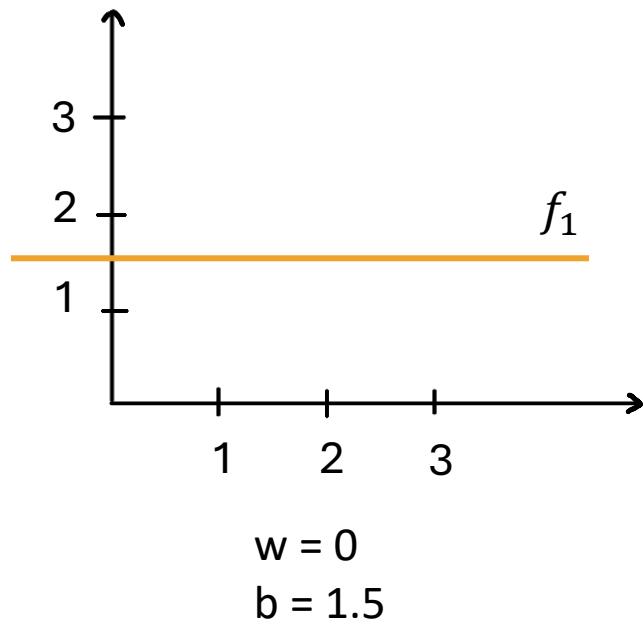
m = number of training examples

$(x^{(i)}, y^{(i)})$ = i^{th} training example

Model: $f_{w,b}(x) = wx + b$
 w, b : parameters

How w and b control the model

Model: $f_{w,b}(x) = wx + b$
 w, b : parameters



Cost function

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2 = \frac{1}{m} (L^{(1)} + L^{(2)} + \dots + L^{(m)})$$

m : number of training examples

\hat{y} : model prediction

y : true label of data

J : cost function

L : loss function

Cost function

Model: $f_{w,b}(x) = wx + b$

w, b : parameters

Cost function:

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$

Goal:

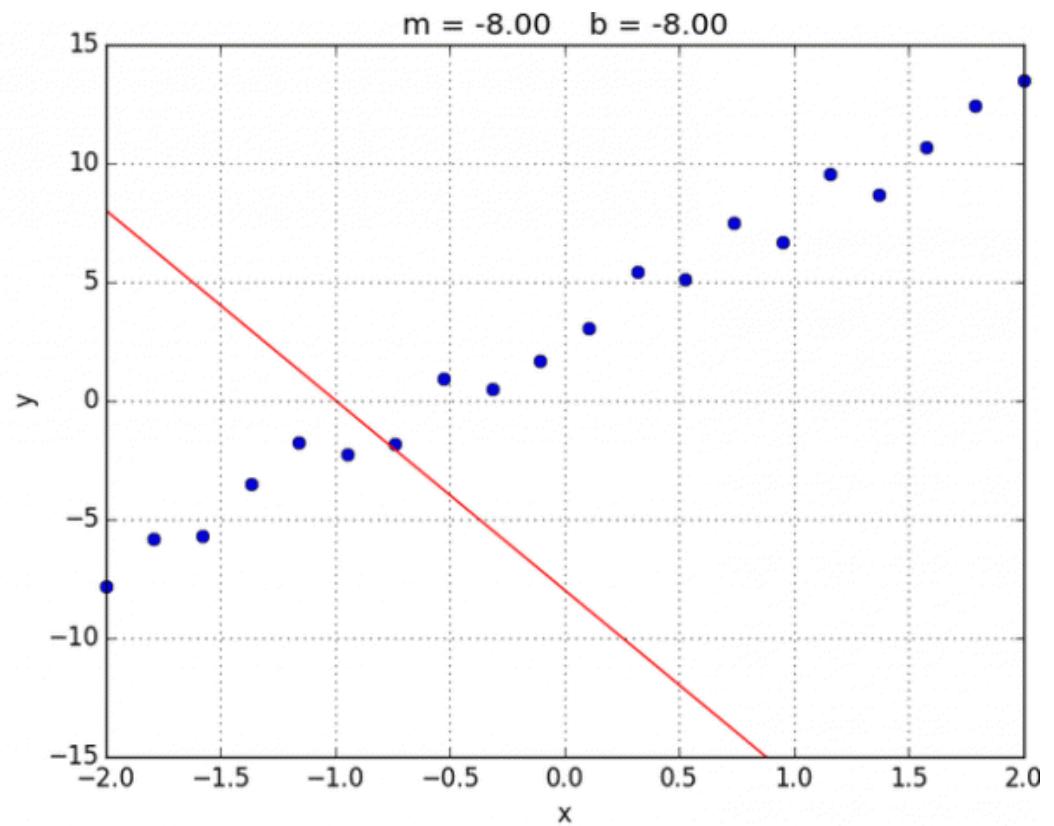
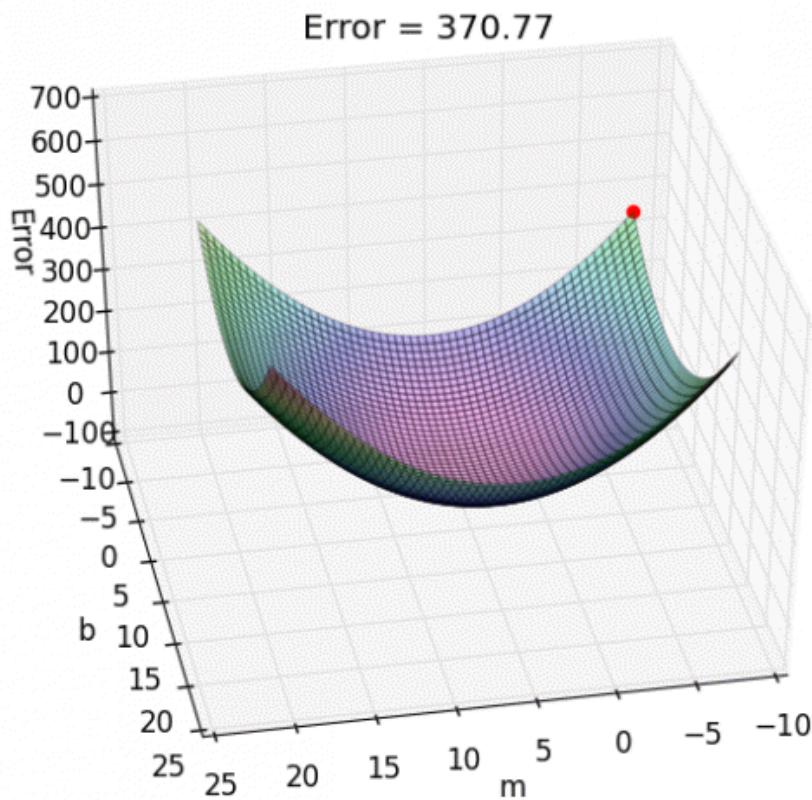
$$\text{Minimize}_{w,b} J(w, b)$$

Quiz

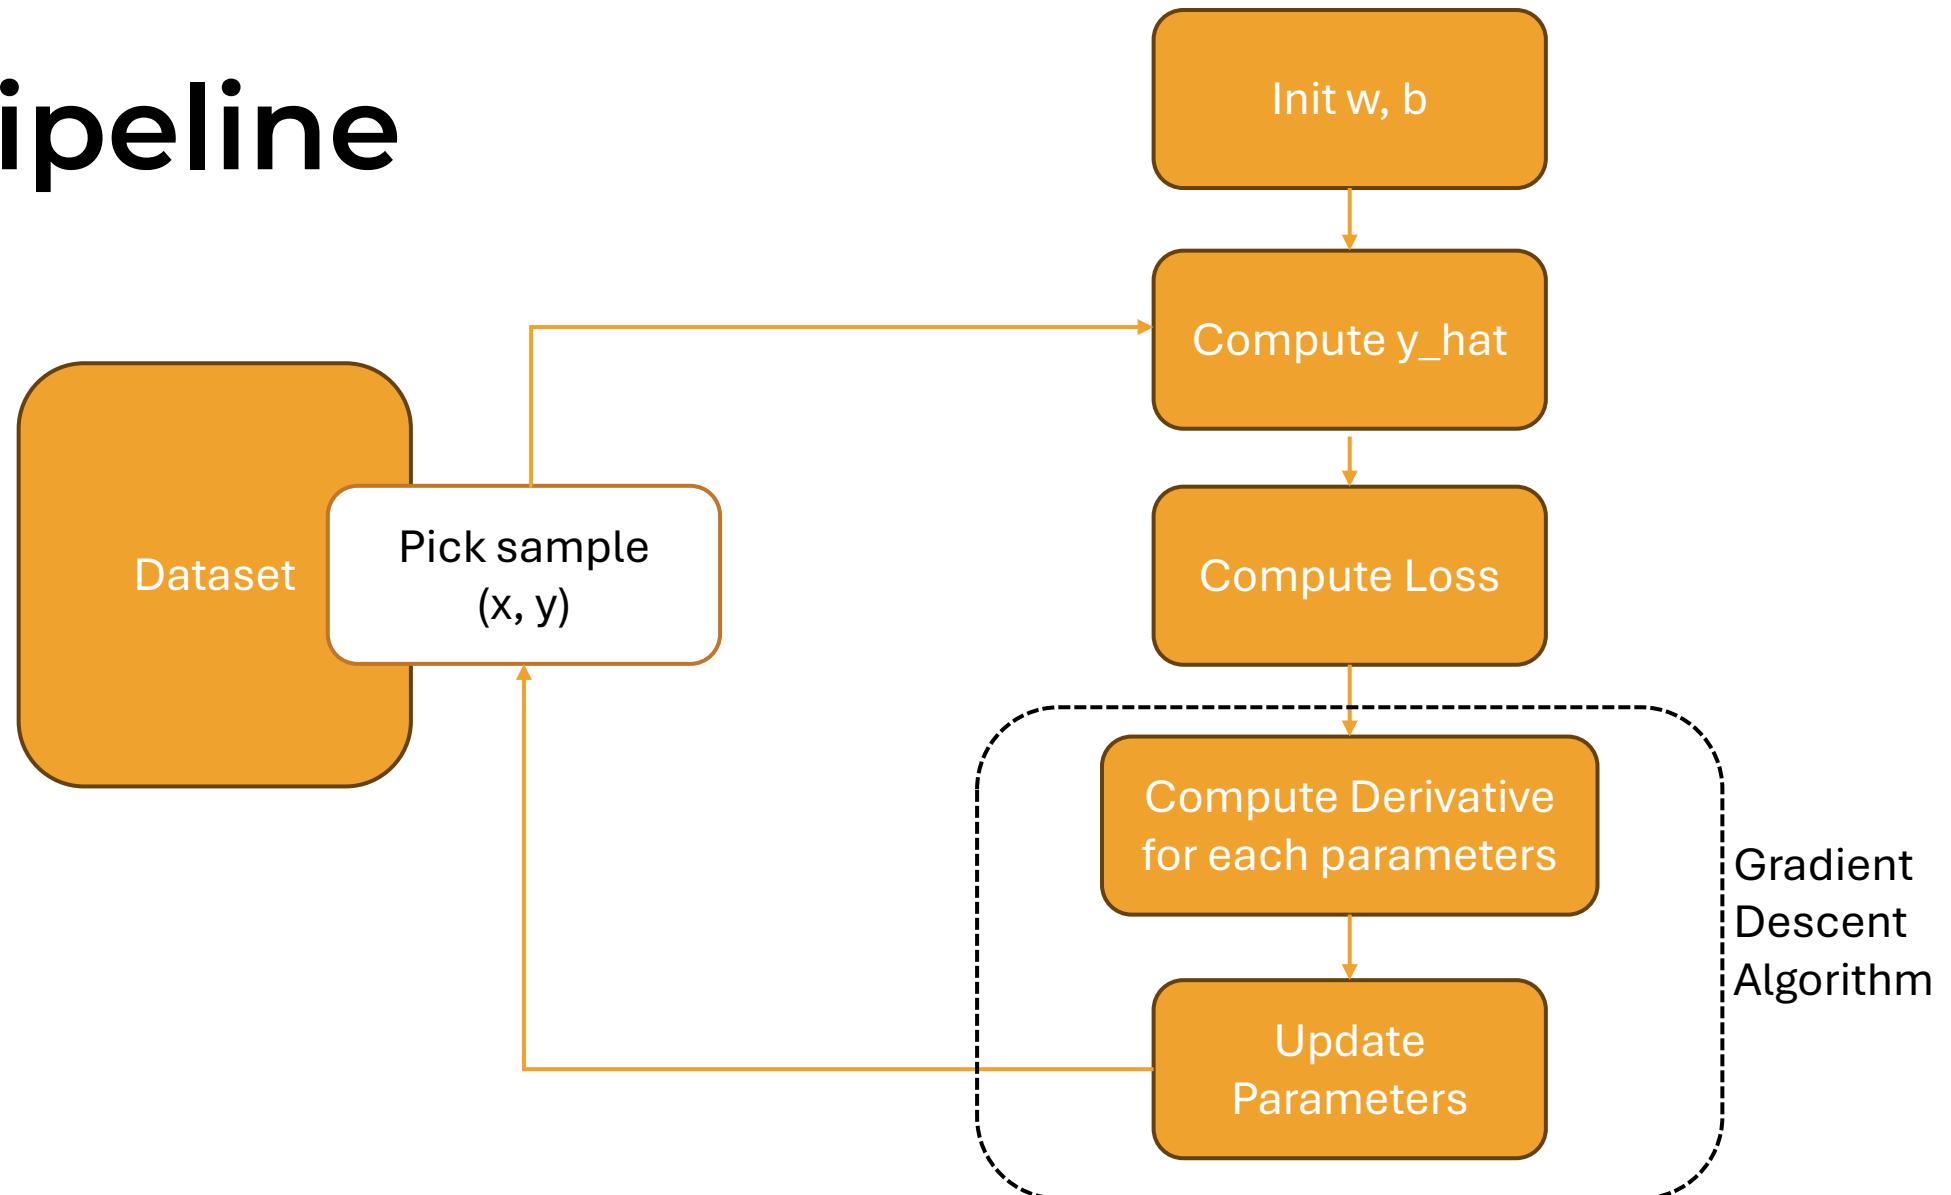
Heigh (cm) x	Weight (kg) y	Predicted $\hat{y} = w * x^{(i)} + b$	Square Error $L = \frac{1}{2}(\hat{y}^{(i)} - y^{(i)})^2$	Loss function $\frac{1}{m} \sum_{i=1}^m L^{(i)}$
147	49	- 2939.5	L1	$= (L^1 + L^2 + \dots + L^8) / m = J$
150	50	..	L2	
153	51	..	L3	
155	52	
158	54	
160	56	
163	58	
165	59	..	L8	

1. What is $x^{(5)}$ and its corresponding y
2. Define m?
3. Assume $w = -20$, $b = 0.5$. Calculate $\hat{y}^{(6)}$
4. Assume $w = -20$, $b = 0.5$. Calculate $J_{w,b}$

Visualize Cost Function



LR pipeline



Gradient Descend

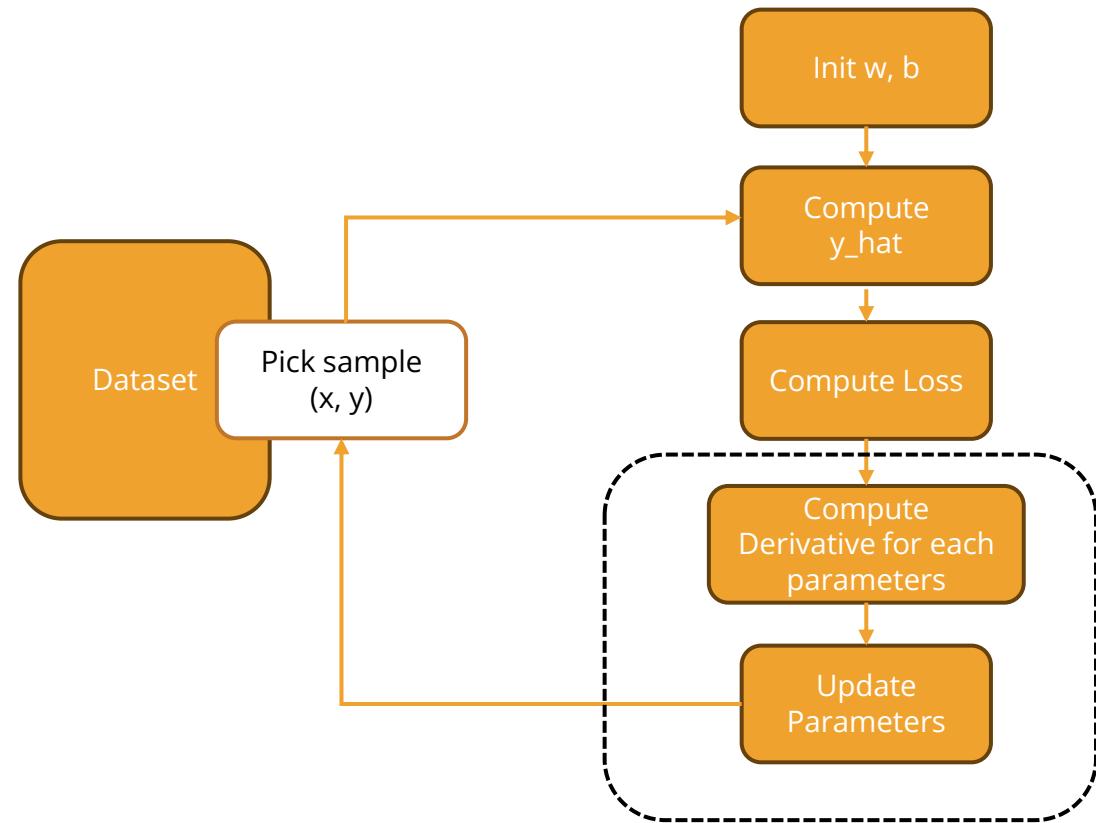
Model: $f_{w,b}(x) = wx + b$
 w, b : parameters

Mean Square Error function:

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$

Goal:

$\text{Minimize}_{w,b} J(w, b)$



Outline:

- Start with some w, b (e.g. $w=0, b=0$)
- Keep changing w, b to reduce $J(w, b)$ until we settle or near a minimum

Gradient Descent Algorithm

Repeat these step until convergence:

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$\frac{\partial J(w, b)}{\partial w} = \frac{1}{m} \left(\frac{\partial L^{(1)}}{\partial w} + \frac{\partial L^{(2)}}{\partial w} + \dots + \frac{\partial L^{(m)}}{\partial w} \right)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

$$\frac{\partial J(w, b)}{\partial b} = \frac{1}{m} \left(\frac{\partial L^{(1)}}{\partial b} + \frac{\partial L^{(2)}}{\partial b} + \dots + \frac{\partial L^{(m)}}{\partial b} \right)$$

α : learning rate

Gradient Descent Algorithm

Correct: **Simultaneous** update

$$\text{temp_w} = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$\text{temp_b} = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

$$w = \text{temp_w}$$

$$b = \text{temp_b}$$

Incorrect

$$\text{temp_w} = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$w = \text{temp_w}$$

$$\text{temp_b} = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

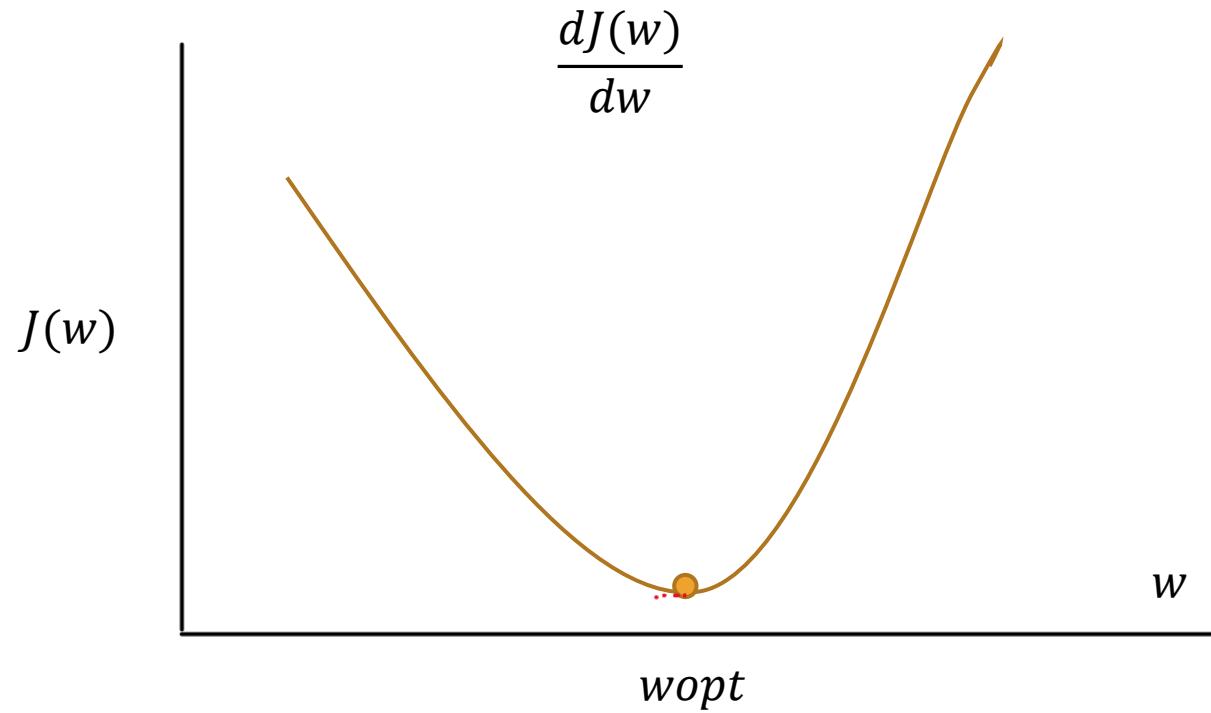
$$b = \text{temp_b}$$

Visualize Cost Function

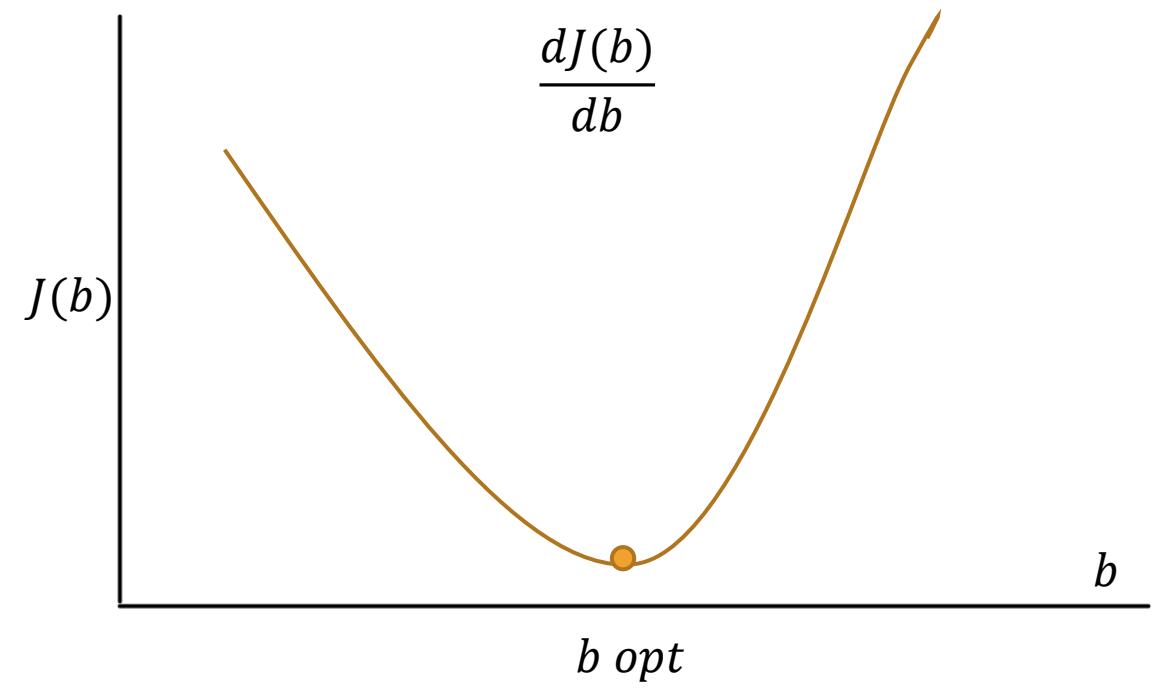
$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

Assume $b = 0$, and different w



Assume $w = 0$, and different b



GD for linear regression

Repeat these step until convergence:

Model: $f_{w,b}(x) = wx + b$

w, b : parameters

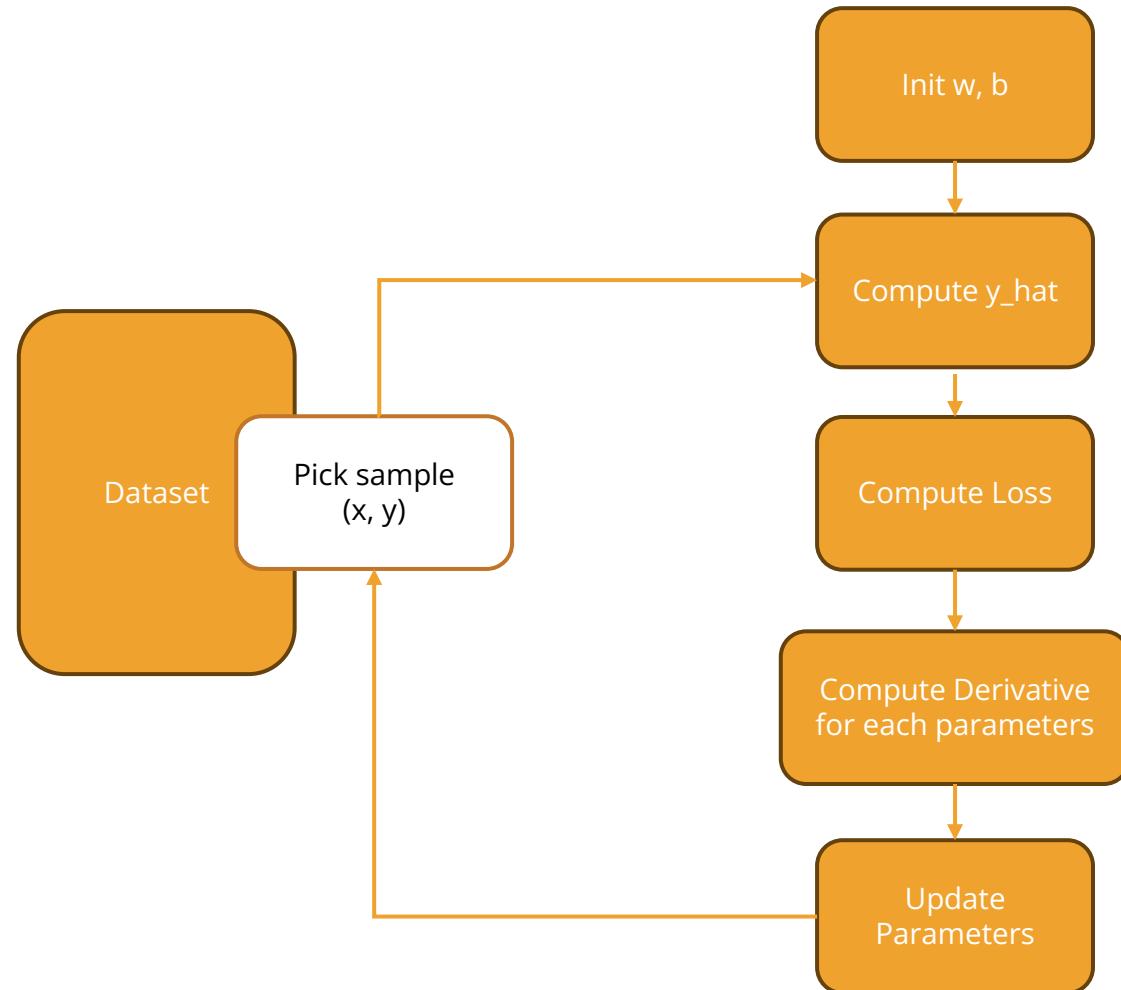
Loss function:

$$L(w, b) = \frac{1}{2} (\hat{y} - y)^2$$

Update parameters:

$$w_{new} = w_{old} - \alpha \frac{\partial}{\partial w} J(w_{old}, b_{old})$$

$$b_{new} = b_{old} - \alpha \frac{\partial}{\partial b} J(w_{old}, b_{old})$$



Batch GD

$$w = w - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x^{(i)}$$

$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})$$

Stochastic GD

$$w = w - 2\alpha(\hat{y}^{(i)} - y^{(i)}) x^{(i)}$$

$$b = b - 2\alpha(\hat{y}^{(i)} - y^{(i)})$$

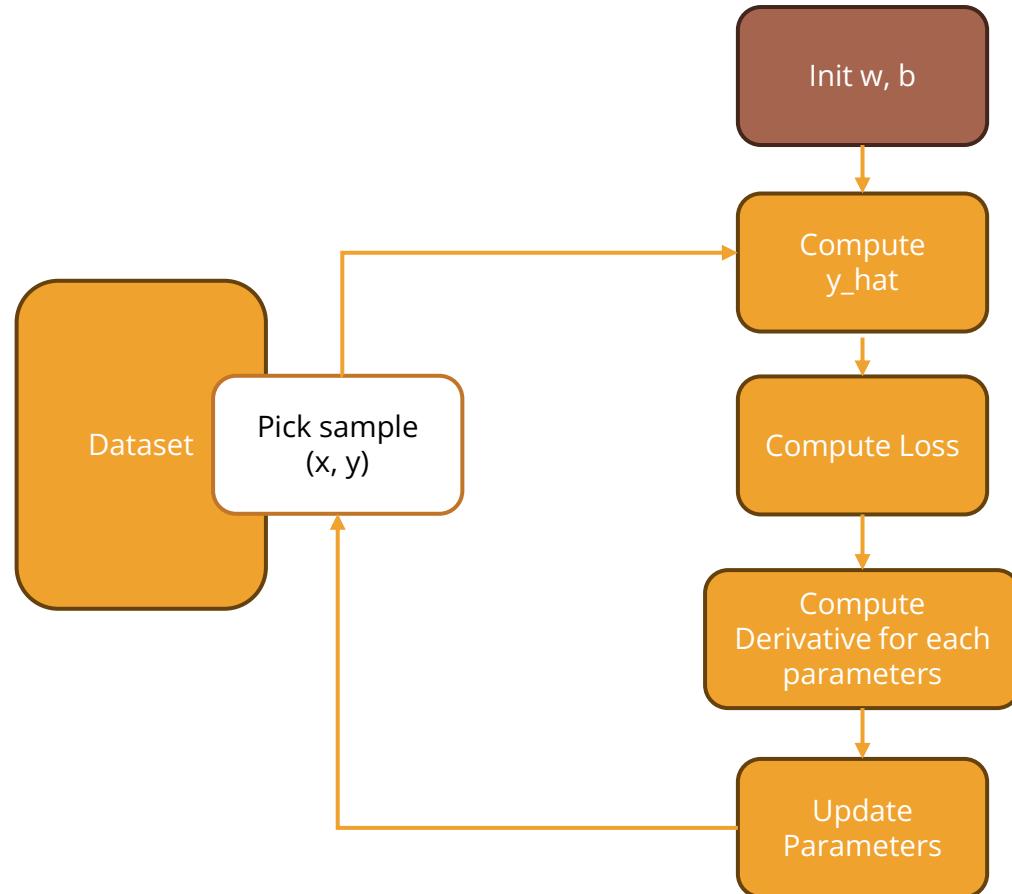
Quiz

Heigh (cm) x	Weight (kg) y	Predicted y $\hat{y}^{(i)} = w * x^{(i)} + b$	Square Error $L = \frac{1}{2}(\hat{y} - y)^2$	Loss function $\frac{1}{m} \sum_{i=1}^m L^{(i)}$
147	50
150	51			
151	51			
155	53			
158	54			
160	56			
163	58			
165	59			

1. Assume $w = 0, b = 0, \alpha = 0.00002$
2. $J(w,b) = ?$
3. Perform SGD for 1 iteration. Calculate `updated_w`, `updated_b`, $J(\text{update}_w, \text{update}_b)$
4. Perform BGD for 1 iteration. Calculate `updated_w`, `updated_b`, $J(\text{update}_w, \text{update}_b)$

Coding exercise

Step 1: Initialize parameters

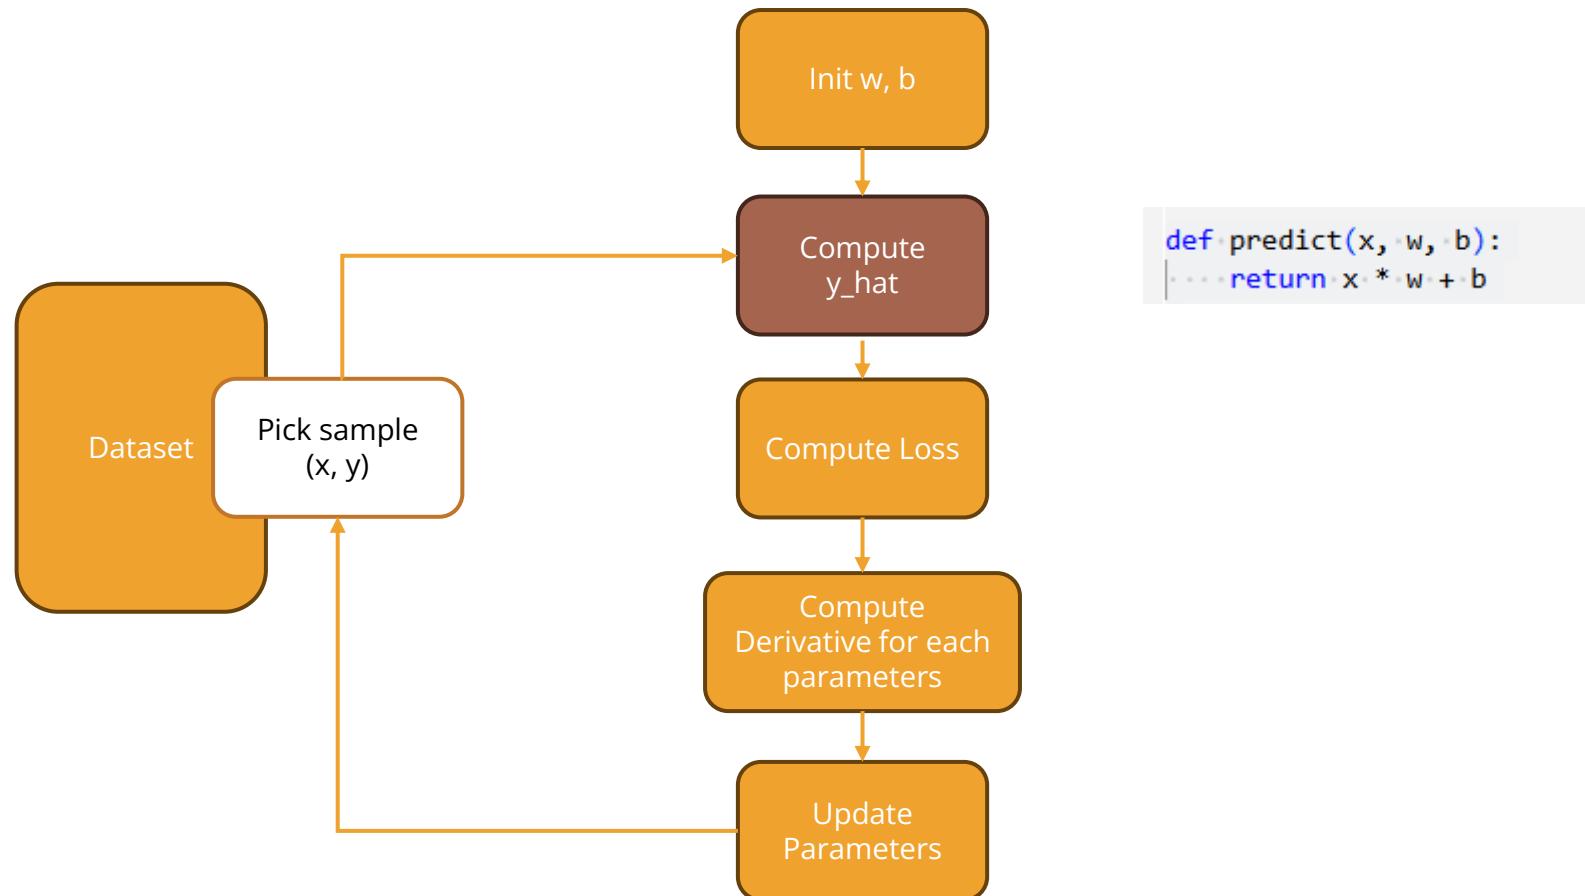


```
import pandas as pd
import numpy as np

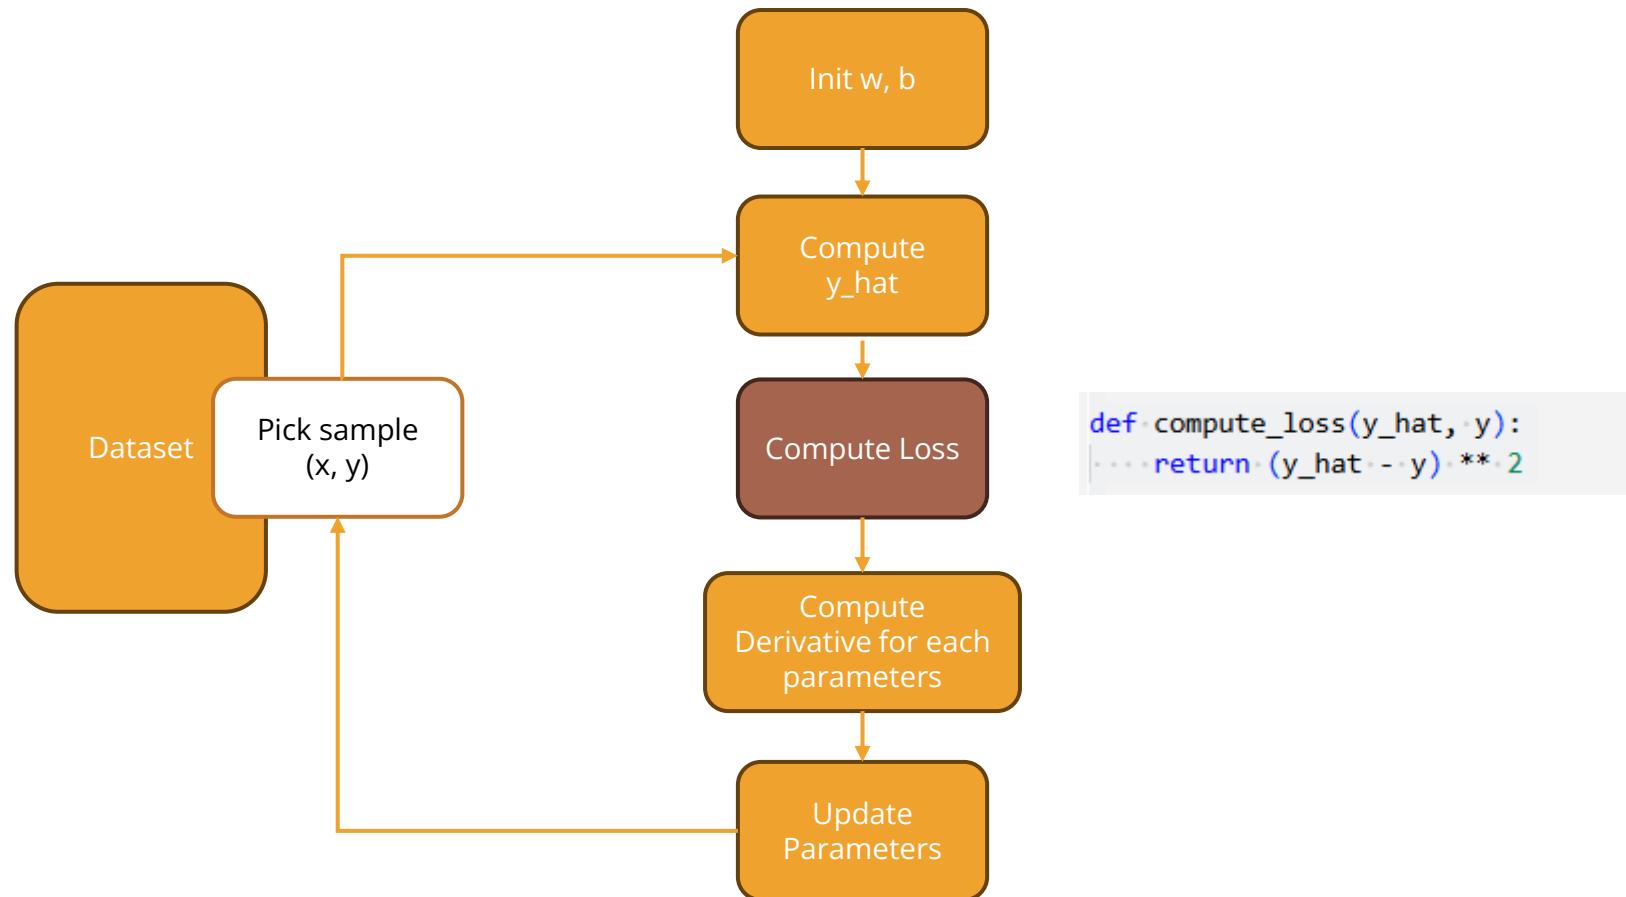
x = pd.read_csv('data.csv')['x'].values
y = pd.read_csv('data.csv')['y'].values

b = 0
w = 0
lr = 0.00002
```

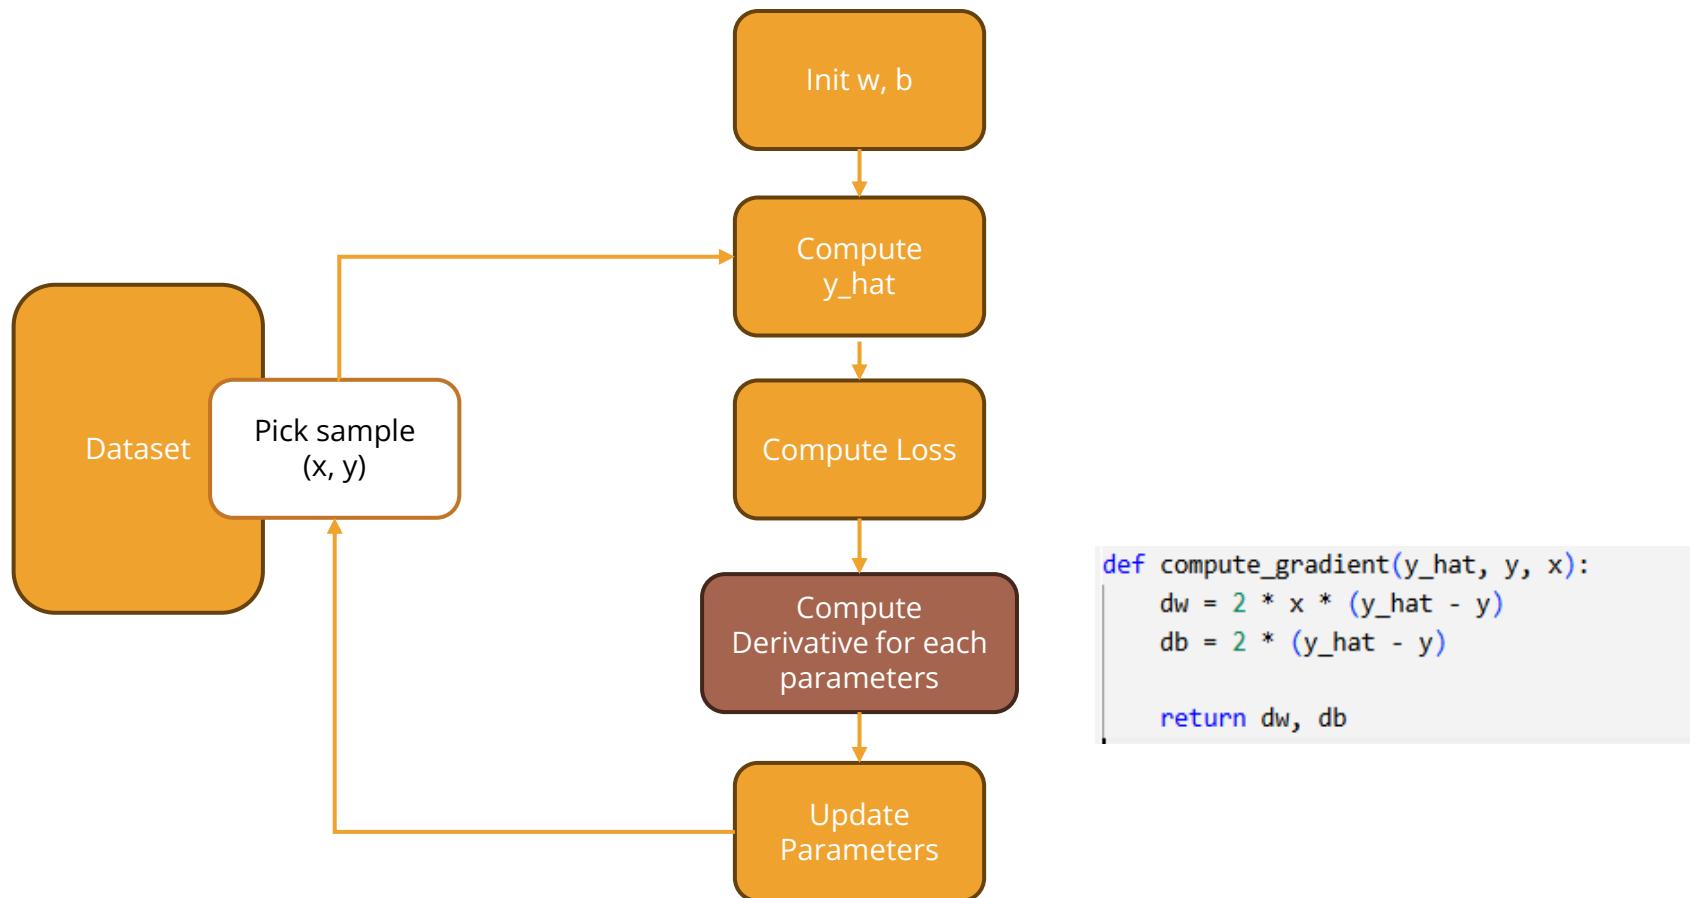
Step 2: Calculate predict



Step 3: Calculate loss

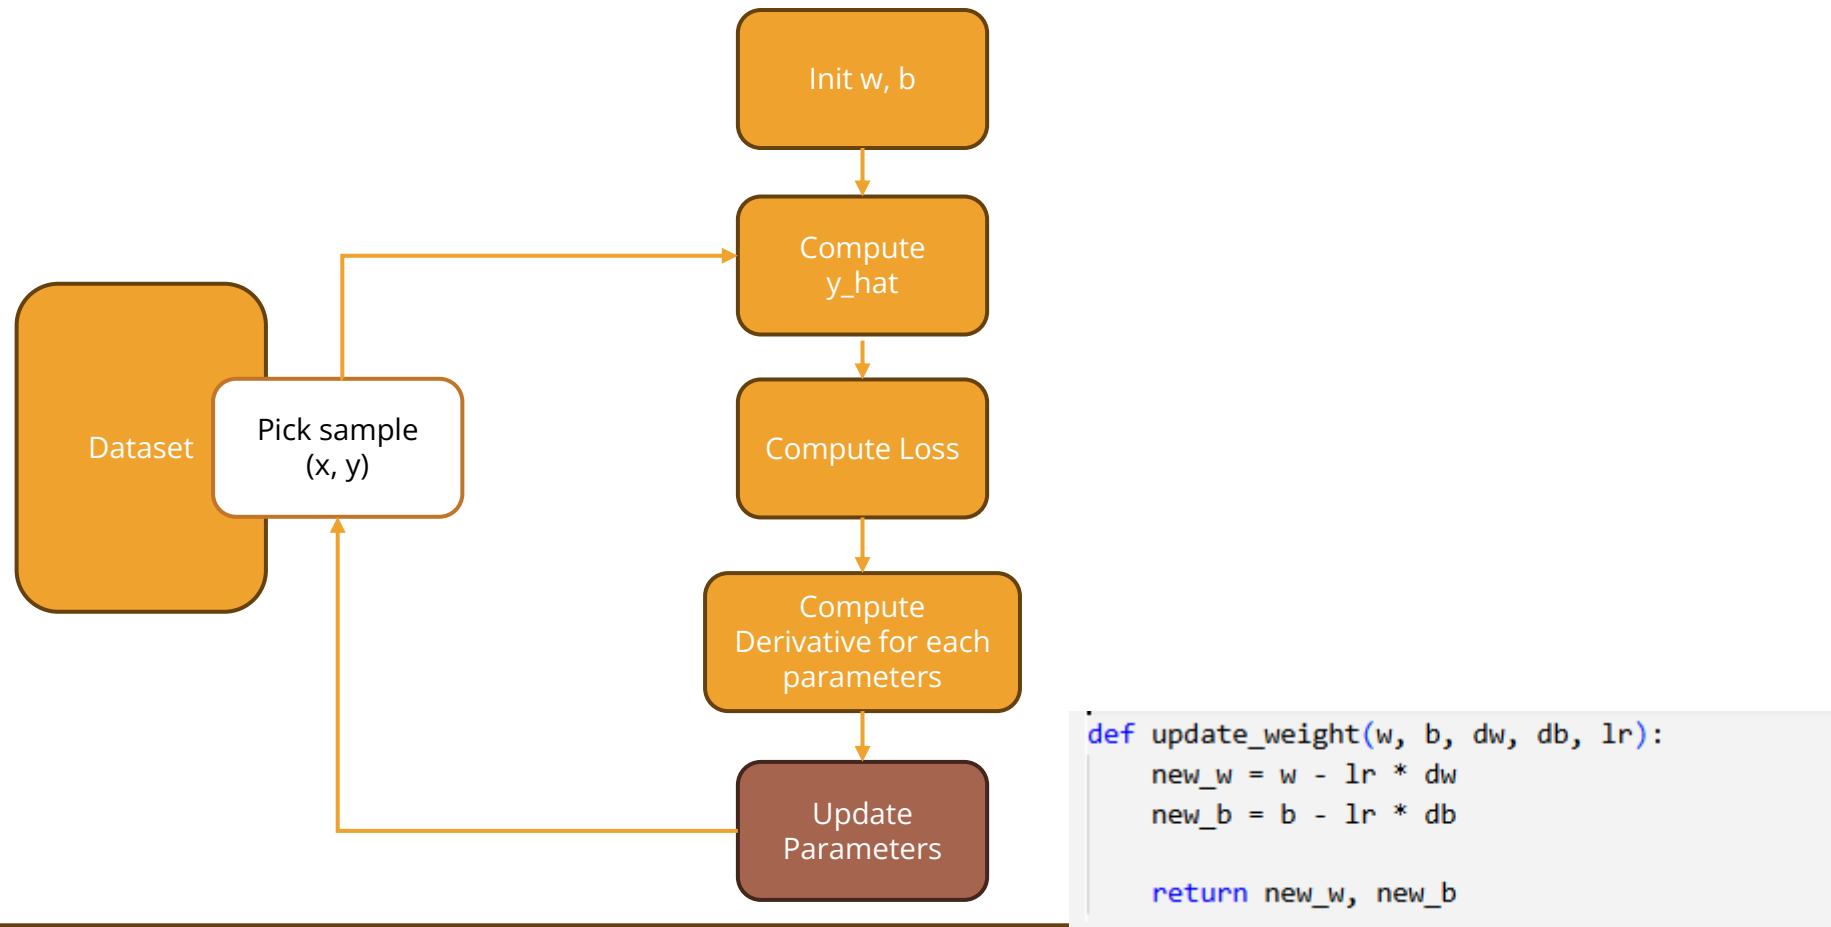


Step 4: Compute gradient



```
def compute_gradient(y_hat, y, x):  
    dw = 2 * x * (y_hat - y)  
    db = 2 * (y_hat - y)  
  
    return dw, db
```

Step 5: Update parameters



THANK YOU